

# Bayesian Optimization under Heavy-tailed Payoffs


Sayak Ray Chowdhury  
(Joint work with Aditya Gopalan)

Department of Electrical Communication Engineering  
Indian Institute of Science

September 25, 2019

# Black-box optimization


Sequentially Maximize  $f : D \rightarrow \mathbb{R}$



►  $f$  unknown,  $D \subset \mathbb{R}^d$

# Black-box optimization

Sequentially Maximize  $f : D \rightarrow \mathbb{R}$




►  $f$  unknown,  $D \subset \mathbb{R}^d$

►  $x^* \in \operatorname{argmax}_{x \in D} f(x)$

# Black-box optimization


Sequentially Maximize  $f : D \rightarrow \mathbb{R}$



- ▶  $f$  unknown,  $D \subset \mathbb{R}^d$
- ▶  $x^* \in \operatorname{argmax}_{x \in D} f(x)$
- ▶ At each round  $t$ :
  - ▶ Learner chooses  $x_t \in D$  based on past
  - ▶ Observes noisy reward  $y_t = f(x_t) + \varepsilon_t$

# Black-box optimization

Sequentially Maximize  $f : D \rightarrow \mathbb{R}$



►  $f$  unknown,  $D \subset \mathbb{R}^d$

►  $x^* \in \operatorname{argmax}_{x \in D} f(x)$

► At each round  $t$ :

- Learner chooses  $x_t \in D$  based on past
- Observes noisy reward  $y_t = f(x_t) + \varepsilon_t$

## Goal

► Minimize Cumulative Regret:  $\sum_{t=1}^T (f(x^*) - f(x_t))$

# Application: Hyperparameter tuning in ML models


- ▶ Hyperparameters in DeepNN training:
  - ▶ Learning rate
  - ▶ Regularizer
  - ▶ Number of hidden layers
  - ▶ Number of units in each layer
  - ▶ Optimizer (SGD, Adagrad, Adam, ...)
  - ▶ Nonlinearity (Relu, Softmax, ...)
  - ▶ ...
- ▶ Black-box optimization:
  - ▶  $D$  : all possible hyperparameter configurations
  - ▶  $f(x)$  : training error for configuration  $x$
  - ▶  $x^*$  : the best hyperparameter

# Traditional approaches

- ▶ Grid search, Random search
  - ▶ Doesn't use information from previous searches
  - ▶ Not good when training time is high
- ▶ Huge (possibly infinite) set of hyperparameters ( $D \subset \mathbb{R}^d$ )
- ▶ Need to make an **educated decision** about where to search
  - ▶ Bayesian optimization

# Bayesian optimization

Sequentially Maximize  $f : D \rightarrow \mathbb{R}$



At each round  $t$ :

- ▶ Learner chooses  $x_t \in D$  based on past knowledge
- ▶ Observes noisy reward  $y_t = f(x_t) + \varepsilon_t$

Existing works: Rewards are light-tailed, e.g., Gaussian or sub-Gaussian (Srinivas et al., 2010, Chowdhury and Gopalan, 2017,...)



# Motivation

Many real life environments exhibit **heavy-tailed** behavior

- ▶ Distribution of delays in data networks
- ▶ Bursty traffic flow distributions
- ▶ Price fluctuations in finance and insurance data

Can we develop **efficient** Bayesian optimization algorithms under **heavy-tailed** environments ?


- ▶ **Efficient algorithm**: Sublinear growth of cumulative regret with time

# Assumptions

- ▶ **Smoothness:**  $f$  lies in **R**eproducing **K**ernel **H**ilbert **S**pace (**RKHS**) of functions  $D \rightarrow \mathbb{R}$ 
  - ▶ **Reproducing property:**  $f(x) = \langle f, k(x, \cdot) \rangle_k$
  - ▶ Induces **smoothness:**  $|f(x) - f(y)| \leq \|f\|_k \|k(x, \cdot) - k(y, \cdot)\|_k$
- ▶ Example kernel
  - ▶ **RBF** kernel:  $k(x, y) = \exp\left(\frac{-\|x-y\|_2^2}{2l^2}\right)$
- ▶ **Heavy-tailed** noise  $\varepsilon_t$  is zero mean and have bounded  $(1 + \alpha)$ -th moment for  $\alpha \in (0, 1]$ .
  - ▶ **Student's- $t$**  distribution with 3 d.o.f. has **bounded variance**

# Algorithm design


**Key idea:** Represent **uncertainty over  $f$**  using **Gaussian Process (GP)**



- Assume **Gaussian Process Prior**  
 $\text{GP}(0, k(x, y))$

# Algorithm design

**Key idea:** Represent **uncertainty over  $f$**  using **Gaussian Process (GP)**



- ▶ Assume **Gaussian Process Prior**  
 $\text{GP}(0, k(x, y))$


- ▶ **Truncate high rewards:**

$$\hat{y}_t = \mathbf{1}_{|y_t| \leq b_t}$$

- ▶  $b_t$  governs the truncation

# Algorithm design

**Key idea:** Represent **uncertainty over  $f$**  using **Gaussian Process (GP)**



► Assume **Gaussian Process Prior**  
 $\text{GP}(0, k(x, y))$

► **Truncate high rewards:**

$$\hat{y}_t = \mathbf{1}_{|y_t| \leq b_t}$$

►  $b_t$  governs the truncation


(Approximate) GP posterior:

$$\mu_t(x) = k_t(x)^T (K_t + \lambda I)^{-1} [\hat{y}_1, \dots, \hat{y}_t]^T$$

$$\sigma_t^2(x) = k(x, x) - k_t(x)^T (K_t + \lambda I)^{-1} k_t(x)$$


# Algorithm 1: Truncated GP-UCB (TGP-UCB)

**Key Idea:** Play the arm with highest Upper Confidence Bound (UCB)



# Algorithm 1: Truncated GP-UCB (TGP-UCB)

**Key Idea:** Play the arm with highest Upper Confidence Bound (UCB)




At each round  $t$ , play:

$$x_t = \operatorname{argmax}_{x \in D} \mu_t(x) + \beta_t \sigma_t(x)$$

# Algorithm 1: Truncated GP-UCB (TGP-UCB)

**Key Idea:** Play the arm with highest **Upper Confidence Bound (UCB)**



At each round  $t$ , play:

$$x_t = \operatorname{argmax}_{x \in D} \mu_t(x) + \beta_t \sigma_t(x)$$

- ▶  $\beta_t$  trades off b/w **exploration** and **exploitation**
- ▶ Well known algorithm under **Gaussian payoffs** (but no reward truncation)



# Cumulative regret of TGP-UCB

## Upper bound of TGP-UCB (Informal)

Cumulative regret of **TGP-UCB** is  $O\left(\gamma_T T^{\frac{2+\alpha}{2(1+\alpha)}}\right)$  with high probability

- ▶  $\gamma_T$  is a function of the kernels and quantifies Reduction in uncertainty about  $f$  after observing rewards
  - ▶ RBF kernel:  $\gamma_T = \text{polylog}(T)$
- ▶ (Recall) assumption: rewards have bounded  $(1 + \alpha)$ -th moment
  - ▶ Bounded variance ( $\alpha = 1$ ): Cumulative regret is  $\tilde{O}(T^{3/4})$
  - ▶ Sublinear growth with  $T$ : TGP-UCB is efficient

# A fundamental lower bound

## Lower bound (Informal)

**RBF kernel:** Expected cumulative regret of any algorithm is  $\Omega(T^{\frac{1}{1+\alpha}})$

- ▶ Bounded variance ( $\alpha = 1$ ): Lower bound is  $\Omega(T^{1/2})$
- ▶ TGP-UCB is efficient, but may be **suboptimal**
- ▶ *Fails to completely nullify the effect of heavy-tail fluctuations*
- ▶ Similar result for **Matérn** kernel

# An (almost) optimal algorithm

## Adaptively Truncated Approximate GP-UCB algorithm

- ▶ Perform **truncation in the feature space** as opposed to truncating raw observations
- ▶ **Challenge:** Feature space is (possibly) **infinite dimensional**
- ▶ **Solution:** Find a “good” **finite-dimensional approximation** of the features
  - ▶ Keep error of approximation in control
- ▶ Perform **feature adaptive truncation** in this approximate feature space

Regret upper bound (Informal):  $O(\gamma_T T^{\frac{1}{1+\alpha}})$

## Algorithm 2: ATA-GP-UCB

$$\mathbf{x}_t = \operatorname{argmax}_{\mathbf{x} \in D} \mu_t(\mathbf{x}) + \beta_t \sigma_t(\mathbf{x})$$

- ▶ Feature approximation of kernel  $\longrightarrow \varphi_t$  (of dimension  $m_t$ )
- ▶  $\Psi_t = [\varphi_t(x_1), \dots, \varphi_t(x_t)]^T \quad V_t = \Psi_t^T \Psi_t + \lambda I_{m_t},$
- ▶ Rows of  $V_t^{-1/2} \Psi_t^T \longrightarrow u_1, \dots, u_{m_t}$
- ▶ Feature adaptive truncation of rewards in each dimension:

$$r_i = \sum_{\tau=1}^t u_{i,\tau} y_\tau \mathbf{1}_{|u_{i,\tau} y_\tau| \leq b_t}$$

### Approximate GP posterior

$$\begin{aligned} \mu_t(x) &= \varphi_t(x)^T V_t^{-1/2} [r_1, \dots, r_{m_t}]^T \\ \sigma_t^2(x) &= k(x, x) - \varphi_t(x)^T \varphi_t(x) + \lambda \varphi_t(x)^T V_t^{-1} \varphi_t(x) \end{aligned}$$

# Fourier features approximation

- ▶ **RBF** kernel:  $k(x, y) = e^{-\frac{(x-y)^2}{2l^2}}$
- ▶ **Random Fourier features** (Rahimi and Recht, 2008):
  - ▶  $\varphi(x) = \frac{1}{\sqrt{m}} [\cos(\omega_1 x), \dots, \cos(\omega_m x), \sin(\omega_1 x), \dots, \sin(\omega_m x)]^T$
  - ▶  $\omega_i \stackrel{\text{i.i.d.}}{\sim} \frac{l}{\sqrt{2\pi}} e^{-\frac{l^2 \omega^2}{2}}$  [Classical method; not useful here]
- ▶ **Quadrature Fourier features** (Mutny and Krause, 2018):
  - ▶  $\varphi(x)_i = \begin{cases} \sqrt{\nu(\omega_i)} \cos\left(\frac{\sqrt{2}}{l} \omega_i x\right) & \text{if } 1 \leq i \leq m, \\ \sqrt{\nu(\omega_{i-m})} \sin\left(\frac{\sqrt{2}}{l} \omega_{i-m} x\right) & \text{if } m+1 \leq i \leq 2m \end{cases}$
  - ▶  $\omega_1, \dots, \omega_m$  : roots of the  $m$ -th **Hermite polynomial**  $H_m$
  - ▶  $\nu(z) = \frac{2^{m-1} m!}{m^2 H_{m-1}(z)^2}$  [Used in this work]


# Nyström approximation

- ▶ **Data dependent:** Approximate the gram matrix  $K_t$
- ▶ Sample  $m_t$  points from  $\{x_1, \dots, x_t\}$  to construct a **dictionary**  $\mathcal{D}_t$
- ▶ Include  $x$  in  $\mathcal{D}_t$  independently with probability  $\sigma_t^2(x)$
- ▶ **Feature approximation:**  $\varphi_t(x) = \left(K_{\mathcal{D}_t}^{1/2}\right)^\dagger k_{\mathcal{D}_t}(x)$ 
  - ▶  $K_{\mathcal{D}_t} = [k(x, y)]_{x, y \in \mathcal{D}_t}$
  - ▶  $k_{\mathcal{D}_t}(x) = [k(x_1, x), \dots, k(x_m, x)]_{x_i \in \mathcal{D}_t}^T$

(Alaoui and Mahony, 2015, Calandriello et. al, 2019)


# Numerical Results

$f$  sampled from RKHS  
(RBF kernel + Impulse noise)



- Reward truncation is necessary to find good estimates


Financial data  
(End of day stock prices)




- Feature adaptive truncation nullify heavy tail fluctuations

# Numerical Results

$f$  sampled from RKHS  
(RBF kernel + Student's- $t$  noise)



$f$  sampled from RKHS  
(RBF kernel + Pareto noise)



- ▶ (Optimal) ATA-GP-UCB is better than (suboptimal) TGP-UCB
- ▶ Nyström embedding is better than Quadrature Fourier features



Thank You